

Assignment No. 5

Problem Statement: Develop and evaluate the K-Nearest Neighbors (KNN) algorithm for both classification and regression tasks..

Objective: Gain a deep understanding of KNN, implement it effectively, and analyze how different parameters influence its performance in classification and regression.

Prerequisite :

1. A Python setup with essential libraries like NumPy, Pandas, Matplotlib, Seaborn, and Scikit-learn.
2. Internet access (for dataset retrieval if necessary).
3. Fundamental knowledge of machine learning and the KNN algorithm.

Theory :

K-Nearest Neighbors (KNN) Algorithm

KNN is a supervised learning technique used for classification and regression. It predicts outcomes based on the similarity between a test data point and its closest training examples.

How KNN Works

1. Choose the number of neighbors (K).
2. Compute the distance (e.g., Euclidean, Manhattan) between the test instance and all training data.
3. Identify the K nearest neighbors.
4. For classification: Assign the most frequent class among the neighbors.
5. For regression: Compute the average (or weighted average) of the neighbors' values.

Choosing the Right K Value

- Small K (e.g., 1) → Sensitive to noise, may lead to overfitting.
- Large K (e.g., 20) → Smoothens decision boundary but may lose local data patterns.

Common Distance Metrics

- Euclidean Distance: Measures direct straight-line distance.
- Manhattan Distance: Calculates distance based on grid-based movements.

Advantages of KNN

Easy to understand and implement.

Works well for small datasets with fewer features.

Disadvantages of KNN

Computationally expensive for large datasets.

Performance is affected by irrelevant features and feature scaling.

Implementation Steps

1. Loading the Dataset
 - Use Pandas to import the dataset.
 - Check dataset properties (.shape, .info(), .isnull().sum()).
2. Data Preprocessing
 - Handle missing values (impute or remove).
 - Encode categorical variables (LabelEncoder, OneHotEncoder).
 - Normalize numerical data (Min-Max Scaling, Standardization).
3. Train-Test Split
 - Use train_test_split to split data (e.g., 80% training, 20% testing).
4. KNN for Classification
 - Use KNeighborsClassifier from `sklearn.neighbors`.
 - Train model and evaluate using accuracy, precision, recall, and confusion matrix.
5. KNN for Regression
 - Use KNeighborsRegressor from `sklearn.neighbors`.
 - Evaluate using Mean Squared Error (MSE) and R-squared Score.
6. Hyperparameter Tuning
 - Experiment with different K values.
 - Compare performance with different distance metrics (Euclidean, Manhattan, Minkowski).
7. Data Visualization
 - Plot decision boundaries for classification.
 - Show the impact of K value on accuracy.
 - Compare actual vs. predicted values in regression.

CODE & OUTPUT :

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, precision_score, recall_score
```

```
[4]: df = pd.read_csv('/Users/pranavashokdivekar/this_mac/Machine Learning/heart.csv')
```

```
[5]: print(df.head(5))
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

```
[6]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    age        1025 non-null    int64
1    sex        1025 non-null    int64
2    cp         1025 non-null    int64
3    trestbps   1025 non-null    int64
4    chol       1025 non-null    int64
5    fbs        1025 non-null    int64
6    restecg    1025 non-null    int64
7    thalach    1025 non-null    int64
8    exang      1025 non-null    int64
9    oldpeak    1025 non-null    float64
10   slope      1025 non-null    int64
11   ca         1025 non-null    int64
12   thal       1025 non-null    int64
13   target     1025 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
None
```

```
[7]: print(df.describe())
```

	age	sex	cp	trestbps	chol	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	54.434146	0.695610	0.942439	131.611707	246.000000	
std	9.072290	0.460373	1.029641	17.516718	51.59251	
min	29.000000	0.000000	0.000000	94.000000	126.000000	
25%	48.000000	0.000000	0.000000	120.000000	211.000000	
50%	56.000000	1.000000	1.000000	130.000000	240.000000	
75%	61.000000	1.000000	2.000000	140.000000	275.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	

	fbs	restecg	thalach	exang	oldpeak	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	0.149268	0.529756	149.114146	0.336585	1.071512	
std	0.356527	0.527878	23.005724	0.472772	1.175053	
min	0.000000	0.000000	71.000000	0.000000	0.000000	
25%	0.000000	0.000000	132.000000	0.000000	0.000000	
50%	0.000000	1.000000	152.000000	0.000000	0.800000	
75%	0.000000	1.000000	166.000000	1.000000	1.800000	
max	1.000000	2.000000	202.000000	1.000000	6.200000	

	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

```
[8]: print(df.isnull().sum())
```

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

```
[10]: # Check if the 'id' column exists in the dataset and drop it (to remove unnecessary identifiers)
      if 'id' in df.columns:
          df.drop(columns=['id'], inplace=True)

      # Check if the 'Unnamed: 32' column exists in the dataset and drop it (likely an extra unnamed column)
      if 'Unnamed: 32' in df.columns:
          df.drop(columns=['Unnamed: 32'], inplace=True)
```

```
• [11]: # Separate the features (X) and target variable (y)
      X = df.iloc[:, 1:] # Select all columns except the first one as features
      y = df.iloc[:, 0]  # Select the first column as the target variable
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=2)
```

```
• [12]: # Create a K-Nearest Neighbors (KNN) classifier with k=3 (3 neighbors)
      knn = KNeighborsClassifier(n_neighbors=3)
      knn.fit(X_train, y_train)
```

```
[12]: KNeighborsClassifier
      KNeighborsClassifier(n_neighbors=3)
```

```
• [13]: # Predict the target values
      y_pred = knn.predict(X_test)
```

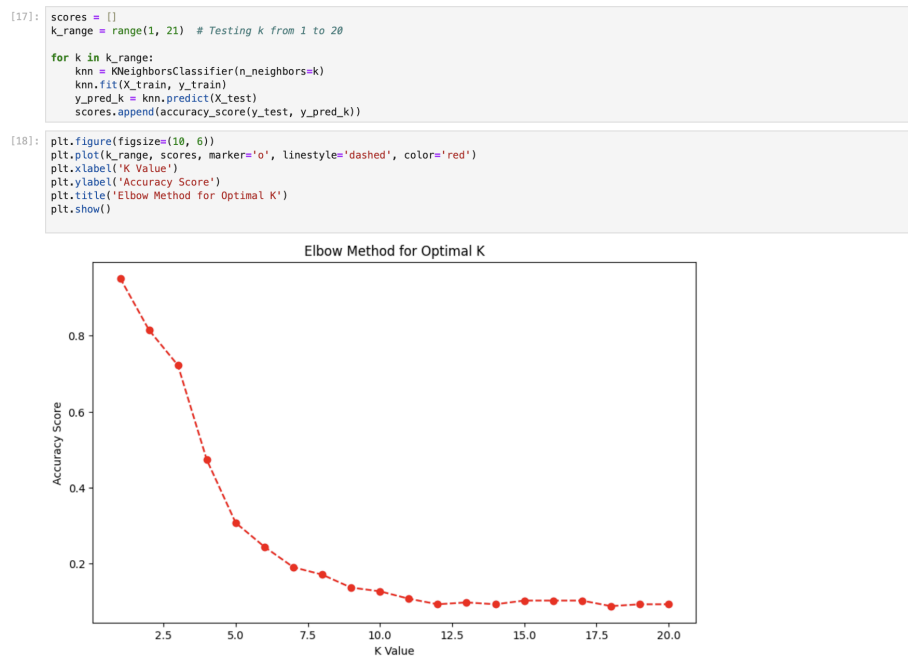
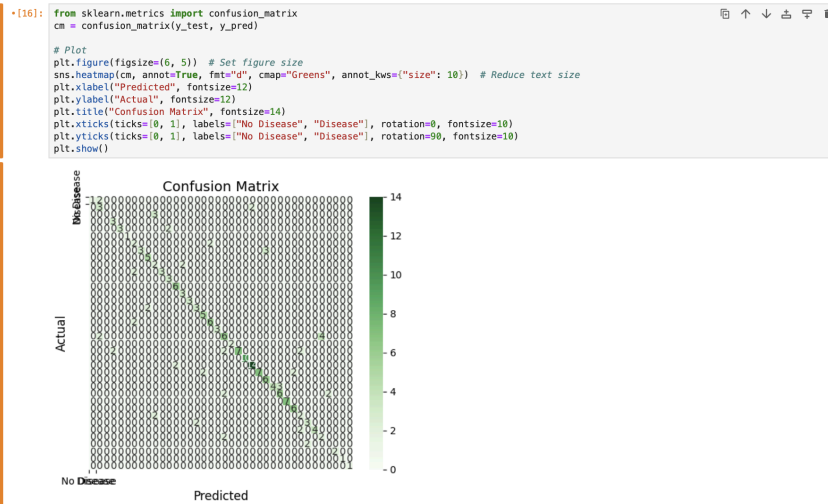
```
• [14]: # Calculate the accuracy
      accuracy = accuracy_score(y_test, y_pred)
      # Calculate the precision (ratio of correctly predicted positive observations to total predicted positives)
      precision = precision_score(y_test, y_pred, average='macro', zero_division=1)
      # Calculate recall
      recall = recall_score(y_test, y_pred, average='macro', zero_division=1)
      # Compute the confusion matrix
      cm = confusion_matrix(y_test, y_pred)
      # Compute specificity
      specificity = cm[0,0] / (cm[0,0] + cm[0,1])
```

```
[15]: print("Accuracy:", accuracy)
      print("Precision:", precision)
      print("Recall:", recall)
      print("Specificity:", specificity)
      print("Confusion Matrix:\n", cm)
      print("Classification Report:\n", classification_report(y_test, y_pred))
```

```
Accuracy: 0.7219512195121951
Precision: 0.7658208020050126
Recall: 0.7360856249014144
Specificity: 0.3333333333333333
Confusion Matrix:
[[1 2 0 ... 0 0 0]
 [0 3 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 2 0 0]
 [0 0 0 ... 0 1 0]
 [0 0 0 ... 0 0 1]]
Classification Report:
              precision    recall  f1-score   support

 34         1.00         0.33         0.50         3
 35         0.43         0.60         0.50         5
 37         0.00         0.00         0.00         3
 38         0.60         1.00         0.75         3
 39         1.00         0.60         0.75         5
 40         1.00         1.00         1.00         1
 41         0.33         0.50         0.40         4
 42         1.00         0.50         0.67         6
 43         0.71         1.00         0.83         5
 44         0.29         0.50         0.36         4
 45         1.00         0.60         0.75         5
 46         0.60         1.00         0.75         3
 47         0.75         1.00         0.86         6
 48         0.60         1.00         0.75         3
 49         1.00         1.00         1.00         3
 50         0.60         0.60         0.60         5
 51         0.71         1.00         0.83         5
 52         0.75         0.75         0.75         8
 53         1.00         1.00         1.00         3
 54         0.50         0.50         0.50         12
 55         1.00         1.00         1.00         2
 56         1.00         0.54         0.70         13
 57         1.00         1.00         1.00         10
 58         0.88         0.88         0.88         16
 59         1.00         0.64         0.78         11
 60         0.67         1.00         0.80         6
 61         1.00         0.57         0.73         7
 62         0.67         0.60         0.63         10
 63         1.00         1.00         1.00         7
 64         0.75         1.00         0.86         6
 65         0.33         0.50         0.40         4
 66         0.60         0.60         0.60         5
 67         1.00         0.67         0.80         6
 68         0.33         0.50         0.40         4
 69         0.00         0.00         0.00         2
 70         1.00         1.00         1.00         2
 76         1.00         1.00         1.00         1
 77         1.00         1.00         1.00         1

 accuracy          0.72         0.72         0.72        205
 macro avg         0.74         0.74         0.71        205
 weighted avg      0.77         0.72         0.72        205
```



Github :- <https://github.com/Pranav-Divekar/Machine-learning->

Conclusion:

The KNN model's performance depended on K value selection and feature scaling. Lower K led to overfitting, while higher K caused underfitting. Proper preprocessing and tuning improved accuracy, making KNN effective for classification and regression while highlighting its strengths and limitations.