

## Assignment No. 4

**Problem Statement:** Understand and implement the Naïve Bayes classification algorithm.

**Objective:**

1. Grasp the Concept – Gain a solid understanding of the mathematical foundation behind the Naïve Bayes algorithm, focusing on Bayes' Theorem and conditional probability principles.
2. Apply Naïve Bayes – Implement the algorithm on a dataset and assess its effectiveness.
3. Assess Performance – Evaluate the classifier's accuracy, precision, recall, and F1-score to determine its efficiency.

**Prerequisite :**

1. Python Environment – Set up a Python workspace with key libraries such as `pandas`, `numpy`, `matplotlib`, `seaborn`, and `scikit-learn`.
2. Foundational Knowledge – Have a basic grasp of Python programming, statistics, and core machine learning concepts.
3. Statistical Understanding – Comprehend probability, conditional probability, and Bayes' Theorem.
4. Machine Learning Basics – Be familiar with classification methods and model evaluation metrics.

**Theory :**

Naïve Bayes is a probabilistic classification algorithm based on **Bayes' Theorem**, used to determine the likelihood of a class given specific features. The term "*naïve*" comes from the assumption that all features are **independent**, which may not always hold true in real-world datasets. Despite this simplification, it remains highly effective for various applications.

Naïve Bayes is built upon **Bayes' Theorem**

Where:

1.  $P(A|B)P(A|B)$  = Probability of event A occurring given that event B has occurred (**posterior probability**).
2.  $P(B|A)P(B|A)$  = Probability of event B occurring given that event A has occurred (**likelihood**).
3.  $P(A)P(A)$  = Prior probability of event A occurring.
4.  $P(B)P(B)$  = Total probability of event B occurring.

**Feature Independence** – Each feature is assumed to contribute independently to the probability of a class label.

**Equal Feature Importance** – All features are considered equally significant in making predictions.

**Conditional Independence** – Given the class label, features are assumed not to be dependent on each other.

## 5. Steps in Naïve Bayes Classification

1. Data Preprocessing – Load the dataset, clean missing values, and prepare the feature-target variables.
2. Compute Prior Probabilities – Calculate  $P(A)$  for each class.
3. Calculate Likelihood – Determine  $P(B | A)$  based on the selected Naïve Bayes model (Gaussian, Multinomial, or Bernoulli).
4. Apply Bayes' Theorem – Compute the posterior probability  $P(A | B)$  for each class and classify based on the highest probability.
5. Evaluate Performance – Assess model accuracy using metrics like precision, recall, F1-score, and overall accuracy.

## Advantages of Naïve Bayes

- Efficient and Fast – Works well with large and high-dimensional datasets.
- Handles Missing Data – Can still perform well even if some features are missing.
- Effective on Small Datasets – Requires less training data than many other classifiers.
- Simple and Interpretable – Easy to understand and implement.
- Excels in Text Classification – Widely used for spam filtering, sentiment analysis, and document categorization.

## Limitations of Naïve Bayes

- Unrealistic Feature Independence Assumption – Often, real-world features are correlated, which can affect accuracy.
- Zero Probability Problem – If a feature value is missing in the training set, the model assigns zero probability to it (solved using Laplace Smoothing).
- Poor Performance with Highly Correlated Features – When features are dependent, the model may produce incorrect classifications.

- Limited for Complex Datasets – Not ideal for problems requiring feature interactions.

## 5. Code & Output

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

• [3]: df = pd.read_csv(/Users/pranavashokdivekar/this_mac/Machine Learning/adult.csv', header=None, sep=',\s')

[5]: df.head()

[5]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

```

• [6]: #Renaming the columns
col_names = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status', 'occupation', 'relationship',
'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country', 'income']
df.columns = col_names
df.columns

[6]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
'marital_status', 'occupation', 'relationship', 'race', 'sex',
'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
'income'],
dtype='object')

• [9]: categorical = [var for var in df.columns if df[var].dtype=='O']

print('There are {} categorical variables\n'.format(len(categorical)))

print('The categorical variables are :\n\n', categorical)

There are 9 categorical variables

The categorical variables are :

['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country', 'income']

```

```

: df[categorical].isnull().sum()

: workclass      0
education        0
marital_status   0
occupation       0
relationship     0
race             0
sex             0
native_country   0
income          0
dtype: int64

```

```

: for var in categorical:
    print(df[var].value_counts())

workclass
Private      22696
Self-emp-not-inc 2541
Local-gov    2093
?            1836
State-gov    1298
Self-emp-inc 1116
Federal-gov  960
Without-pay  14
Never-worked 7
Name: count, dtype: int64
education
HS-grad      10501
Some-college 7291
Bachelors    5355
Masters       1723
Assoc-voc     1382
11th          1175
Assoc-acdm    1067
10th          933
7th-8th       646
Prof-school   576
9th           514
12th          433
...

```

```

• [13]: df.workclass.unique()

[13]: array(['State-gov', 'Self-emp-not-inc', 'Private', 'Federal-gov',
        'Local-gov', '?', 'Self-emp-inc', 'Without-pay', 'Never-worked'],
        dtype=object)

• [14]: df.workclass.value_counts()

[14]: workclass
      Private      22696
      Self-emp-not-inc  2541
      Local-gov      2093
      ?             1836
      State-gov      1298
      Self-emp-inc    1116
      Federal-gov     960
      Without-pay     14
      Never-worked     7
      Name: count, dtype: int64

• [16]: df['workclass'].replace('?', np.nan, inplace=True)

• [17]: df.workclass.value_counts()

[17]: workclass
      Private      22696
      Self-emp-not-inc  2541
      Local-gov      2093
      State-gov      1298
      Self-emp-inc    1116
      Federal-gov     960
      Without-pay     14
      Never-worked     7
      Name: count, dtype: int64

[18]: # check labels in occupation variable
      df.occupation.unique()

[18]: array(['Adm-clerical', 'Exec-managerial', 'Handlers-cleaners',
        'Prof-specialty', 'Other-service', 'Sales', 'Craft-repair',
        'Transport-moving', 'Farming-fishing', 'Machine-op-inspct',
        'Tech-support', '?', 'Protective-serv', 'Armed-Forces',
        'Priv-house-serv'], dtype=object)

```

```

[21]: # replace '?' values in occupation variable with 'NaN'
      df['occupation'].replace('?', np.nan, inplace=True)

[22]: # again check the frequency distribution of values in occupation variable
      df.occupation.value_counts()

[22]: occupation
      Prof-specialty      4140
      Craft-repair      4099
      Exec-managerial    4066
      Adm-clerical      3770
      Sales              3650
      Other-service      3295
      Machine-op-inspct  2002
      Transport-moving   1597
      Handlers-cleaners  1370
      Farming-fishing     994
      Tech-support        928
      Protective-serv     649
      Priv-house-serv     149
      Armed-Forces         9
      Name: count, dtype: int64

[23]: # check labels in native_country variable
      df.native_country.unique()

[23]: array(['United-States', 'Cuba', 'Jamaica', 'India', '?', 'Mexico',
        'South', 'Puerto-Rico', 'Honduras', 'England', 'Canada', 'Germany',
        'Iran', 'Philippines', 'Italy', 'Poland', 'Columbia', 'Cambodia',
        'Thailand', 'Ecuador', 'Laos', 'Taiwan', 'Haiti', 'Portugal',
        'Dominican-Republic', 'El-Salvador', 'France', 'Guatemala',
        'China', 'Japan', 'Yugoslavia', 'Peru',
        'Outlying-US(Guam-USVI-etc)', 'Scotland', 'Trinidad&Tobago',
        'Greece', 'Nicaragua', 'Vietnam', 'Hong', 'Ireland', 'Hungary',
        'Holand-Netherlands'], dtype=object)

```

```
[27]: # again check the frequency distribution of values in native_country variable
```

```
df.native_country.value_counts()
```

```
[27]: native_country
      United-States      29170
      Mexico           643
      Philippines      198
      Germany          137
      Canada           121
      Puerto-Rico       114
      El-Salvador       106
      India            100
      Cuba              95
      England           90
      Jamaica           81
      South             80
      China             75
```

```
[33]: # split X and y into training and testing sets
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
```

```
[34]: # check the shape of X_train and X_test
```

```
X_train.shape, X_test.shape
```

```
[34]: ((22792, 14), (9769, 14))
```

```
[35]: # check data types in X_train
```

```
X_train.dtypes
```

```
[35]: age                int64
workclass            object
fnlwgt              int64
education            object
education_num        int64
marital_status       object
occupation           object
relationship         object
race                object
sex                 object
capital_gain         int64
capital_loss         int64
hours_per_week       int64
native_country       object
dtype: object
```

```
*[36]: categorical = [col for col in X_train.columns if X_train[col].dtypes == 'O']
categorical
```

```
[36]: ['workclass',
      'education',
      'marital_status',
      'occupation',
      'relationship',
      'race',
      'sex',
      'native_country']
```

```
X_train[categorical].isnull().sum()
```

```
[41]: workclass      0
      education    0
      marital_status  0
      occupation   0
      relationship  0
      race         0
      sex          0
      native_country  0
      dtype: int64
```

```
*[42]: X_test[categorical].isnull().sum()
```

```
[42]: workclass      0
      education    0
      marital_status  0
      occupation   0
      relationship  0
      race         0
      sex          0
      native_country  0
      dtype: int64
```

```
[43]: # check missing values in X_train
```

```
X_train.isnull().sum()
```

```
[43]: age                0
workclass            0
fnlwgt              0
education            0
education_num        0
marital_status       0
occupation           0
relationship         0
race                0
sex                 0
capital_gain         0
capital_loss         0
hours_per_week       0
native_country       0
dtype: int64
```

```
[44]: # check missing values in X_test
```

```
X_test.isnull().sum()
```

```
[44]: age          0
      workclass    0
      fnlgt        0
      education    0
      education_num 0
      marital_status 0
      occupation   0
      relationship 0
      race          0
      sex          0
      capital_gain  0
      capital_loss  0
      hours_per_week 0
      native_country 0
      dtype: int64
```

```
[45]: # print categorical variables
      categorical
```

```
[45]: ['workclass',
      'education',
      'marital_status',
      'occupation',
      'relationship',
      'race',
      'sex',
      'native_country']
```

```
[46]: X_train[categorical].head()
```

```
[46]:
```

	workclass	education	marital_status	occupation	relationship	race	sex	native_country
32098	Private	HS-grad	Married-civ-spouse	Craft-repair	Husband	White	Male	United-States
25206	State-gov	HS-grad	Divorced	Adm-clerical	Unmarried	White	Female	United-States
23491	Private	Some-college	Married-civ-spouse	Sales	Husband	White	Male	United-States
12367	Private	HS-grad	Never-married	Craft-repair	Not-in-family	White	Male	Guatemala
7054	Private	7th-8th	Never-married	Craft-repair	Not-in-family	White	Male	Germany

```
[47]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column              Non-Null Count  Dtype
---  --
 0   age                 32561 non-null  int64
 1   workclass           38725 non-null  object
 2   fnlgt               32561 non-null  int64
```

```
[49]: # Modify ColumnTransformer to return a dense array
```

```
preprocessor = ColumnTransformer([
    ('num', StandardScaler(), numerical_features), # Scale numerical features
    ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), categorical_features) # OneHot encode categorical features
])
```

```
[50]: from sklearn.naive_bayes import GaussianNB
      from sklearn.metrics import accuracy_score
```

```
# Convert the sparse matrix to dense array
X_train_dense = X_train.toarray()
X_test_dense = X_test.toarray()

# Train Naive Bayes Model
nb_model = GaussianNB()
nb_model.fit(X_train_dense, y_train)

y_pred_nb = nb_model.predict(X_test_dense)


# Evaluate Naive Bayes
accuracy_nb = accuracy_score(y_test, y_pred_nb)
print(f"Naive Bayes Accuracy: {accuracy_nb:.4f}")
```

```
Naive Bayes Accuracy: 0.5559
```

```
model accuracy score: 0.5559
Training-set accuracy score: 0.5552
Training set score: 0.5552
Test set score: 0.5559
```

```
Class distribution in test set:
income
<=50K    7407
>50K     2362
Name: count, dtype: int64
Null accuracy score: 0.7582
```

```
[51]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1e7b161a420>
```



**Github :-** <https://github.com/Pranav-Divekar/Machine-learning->

**Conclusion:**

Gaussian Naïve Bayes is an effective and simple classification algorithm for numerical datasets. Despite its assumption of feature independence, it performs well in many real-world applications, especially when the data is normally distributed. While it has limitations, such as sensitivity to non-Gaussian distributions and the independence assumption, its efficiency, simplicity, and effectiveness make it a valuable tool in machine learning.